


The learned policy defers, then confirms



Pooled decisions across the paired 24-seed held-out urn eval (seeds 2000-2023, no-announce, q8_0). Base keeps eagerly from the first sighting onward; RL-final passes on the first sighting and keeps once recurrence is confirmed on the second -- the reserve policy's decision rule made visible, not just its aggregate first-sight rate. Positions 3+ have small per-bucket n (3-7 decisions) and are a noisy tail, not part of the claim.